


ECE 463
Introduction to Computer Networks

Lecture: Routing


Sanjay Rao

Router Operation




- When Packet Arrives at Router
 - Examine header to determine intended destination
 - *Look up in table to determine next hop in path*
 - Send packet out appropriate port
- Terminology
 - Each router *forwards* packet to next router
 - Overall goal is to *route* packet from source to destination
- Today's task
 - How to generate the routing table

Routing Hierarchy



- Autonomous System
 - Network administered by single entity
- Intradomain Routing
 - Routing within single AS
 - Typically ~ 200 nodes
- Interdomain Routing
 - Routing between AS's


Graph Model



- Represent each router as node
- Direct link between routers represented by edge
 - Symmetric links \Rightarrow undirected graph
- Edge “cost” $c(x,y)$ denotes measure of difficulty of using link
- Task
 - Determine least cost path from every node to every other node
 - Path cost $d(x,y)$ = sum of link costs

Ways to Compute Shortest Paths

- Centralized
 - Collect graph structure in one place
 - Use standard graph algorithm
 - Disseminate routing tables
- Partially Distributed
 - Every node collects complete graph structure
 - Each computes shortest paths from it
 - Each generates own routing table
 - “Link-state” algorithm
- Fully Distributed
 - No one has copy of graph
 - Nodes construct their own tables iteratively
 - Each sends information about its table to neighbors
 - “Distance-Vector” algorithm



Example Routing Tables

Table for A		
Dest	Cost	Next Hop
A	0	A
B	4	B
C	6	E
D	7	B
E	2	E
F	5	E




Table for B		
Dest	Cost	Next Hop
A	4	A
B	0	B
C	2	F
D	3	F
E	4	F
F	1	F


- Multiple choices for A
 - A-E-F-C-D, A-B-D
- Multiple choices for B as well
 - B-D, B-F-C-D
- A may not know route B takes

Issues

- Ensure packets don't get stuck in a loop
- Find “good” paths
- Adapt to changes in edge costs
- Adapt to failure of nodes


Distance-Vector Method

Initial Table for A		
Dest	Cost	Next Hop
A	0	A
B	4	B
C	∞	—
D	∞	—
E	2	E
F	6	F



- Idea
 - At any time, have cost/next hop of best known path to destination
 - Use cost ∞ when no path known
- Initially
 - Only have entries for directly connected nodes


Distance-Vector Update



- **Update(x,y,z)**
 $d \leftarrow c(x,z) + d(z,y)$ # Cost of path from x to y with first hop z
 if $d < d(x,y)$
 # Found better path
 return d,z # Updated cost / next hop
 else
 return d(x,y), nexthop(x,y) # Existing cost / next hop

Routing Algorithm

- Periodically, every node z sends each neighbor x a copy of its routing table
- When x receives the table, it runs $\text{Update}(x,y,z)$ for every destination y
- Process occurs in “asynchronous” fashion
- Routing tables “eventually” converge



What if Node Fails?

- D & F notice that C isn't responding
- Set entries to ∞
- Iterate




Table for D		
Dst	Cst	Hop
C	∞	—

Table for F		
Dst	Cst	Hop
C	∞	—

Failing Node Iterations

- Stale entry in B propagates to D & F
- What Happened?




Table for B		
Dst	Cst	Hop
C	2	F

**Better
Route**

Table for D		
Dst	Cst	Hop
C	∞	—



Table for D		
Dst	Cst	Hop
C	5	B

**Better
Route**

Table for F		
Dst	Cst	Hop
C	∞	—

Table for F		
Dst	Cst	Hop
C	3	B

Revised Update Rule #1



**Variants: “Split Horizon Rule”,
“Split Horizon with Poison Reverse”**

- **Update(x,y,z)**
 $d \leftarrow c(x,z) + d(z,y)$ # Cost of path from x to y with first hop z
 if $d < d(x,y)$ & $x \neq \text{nexthop}(z,y)$
 # Found better path
 return d,z
 else
 return $d(x,y)$, $\text{nexthop}(x,y)$

Iterations with Revision #2

- Stale entry in B still propagates
- What Happened?




Table for B		
Dst	Cst	Hop
C	2	F

**Better
Route**


Table for D		
Dst	Cst	Hop
C	∞	—

Table for D		
Dst	Cst	Hop
C	5	B

Table for F		
Dst	Cst	Hop
C	∞	—

**No
Change**

Revised Update Rule #2



- $\text{Update}(x,y,z)$
 $d \leftarrow c(x,z) + d(z,y)$ # Cost of path from x to y with first hop z
 if $\text{nexthop}(x,y) = z$ | # Forced update
 $(d < d(x,y) \ \& \ x \neq \text{nexthop}(z,y))$
 # Forced update or found better path
 return d,z
 else
 return $d(x,y), \text{nexthop}(x,y)$

Iterations with Revision #2

Table for B		
Dst	Cst	Hop
C	2	F

Table for D		
Dst	Cst	Hop
C	∞	—

Table for F		
Dst	Cst	Hop
C	∞	—

**Better
Route**

Table for D		
Dst	Cst	Hop
C	5	B


**Forced
Update**

Table for B		
Dst	Cst	Hop
C	∞	—

**Forced
Update**

Table for D		
Dst	Cst	Hop
C	∞	—

**No
Change**



- Forced updates will eliminate false entries
- Forced update rule violates monotonicity
 - Increases $d(x,y)$

Routing Information Protocol (RIP)

- Earliest IP routing protocol (1982 BSD)
 - Ideas in first Arpanet protocols (late 60's)
- Current standard is version 2 (RFC 1723)
- Features
 - Every link has cost 1
 - “Infinity” = 16
 - Limits to networks where everything reachable within 15 hops
- Sending Updates
 - Every router listens for updates on UDP port 520

RIP Updates

- Initial
 - When router first starts, asks for copy of table for every neighbor
 - Uses it to iteratively generate own table
- Periodic
 - Every 30 seconds, router sends copy of its table to each neighbor
 - Neighbors use to iteratively update their tables
- Triggered
 - When entry changes, send copy of entry to neighbors
 - Except for one causing update
 - Neighbors use to update their tables

RIP Staleness / Oscillation Control


- Small Infinity
 - Count to infinity doesn't take very long
- Route Timer
 - Every route has timeout limit of 180 seconds
 - Reached when haven't received update from next hop for 6 periods
 - If not updated, set to infinity
- Behavior
 - When router or link fails, can take minutes to stabilize
 - Lots of subtlety to get good implementation (RFC 1058).

Features of Distributed Algorithms

- Desirable in Network Setting
 - Every node operates in purely local way
 - No central control or global synchronization
 - Only communication between direct neighbors
- Not Difficult to Handle Static System
 - Monotonicity guarantees convergence
- Difficult in Dynamically-Changing System
 - Anything that reduces link cost OK
 - Iterations will converge to reflect reduced costs
 - Anything that increases link cost problematic
 - Iterations will converge, but possibly to wrong values
 - Changing update rule can lead to convergence problems
 - Violate monotonicity

---Link-State Routing---

Graph Model



- Represent each router as node
- Direct link between routers represented by edge
 - Symmetric links \Rightarrow undirected graph
- Edge “cost” $c(x,y)$ denotes measure of difficulty of using link
- Task
 - Determine least cost path from every node to every other node
 - Path cost $d(x,y)$ = sum of link costs

Optimal Routing Table





Table for A			Table for B			Table for C			Table for D			Table for E			Table for F		
Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop
A	0	A	A	3	A	A	5	F	A	6	B	A	2	A	A	4	B
B	3	B	B	0	B	B	2	F	B	3	B	B	4	F	B	1	B
C	5	B	C	2	F	C	0	C	C	1	C	C	4	F	C	1	C
D	6	B	D	3	D	D	1	D	D	0	D	D	5	F	D	2	C
E	2	E	E	4	F	E	4	F	E	5	C	E	0	E	E	3	E
F	4	B	F	1	F	F	1	F	F	2	C	F	3	F	F	0	F

Ways to Compute Shortest Paths

- Centralized
 - Collect graph structure in one place
 - Use standard graph algorithm
 - Disseminate routing tables
- Partially Distributed 
 - Every node collects complete graph structure
 - Each computes shortest paths from it
 - Each generates own routing table
 - “Link-state” algorithm
- Fully Distributed
 - No one has copy of graph
 - Nodes construct their own tables iteratively
 - Each sends information about its table to neighbors
 - “Distance-Vector” algorithm

Why Not Fully Distributed?

- Original RIP Protocol
 - Link cost = 1
 - Treat 16 like infinity
- Good Features
 - Simple to implement
 - Completed distributed
- Bad Features
 - Potentially slow convergence
 - Count-to-infinity
 - Tables could have loops when not yet converged
 - Had to use very weak metrics & features
 - Restricted to spanning-tree routes

Link State Protocol Concept

- Every Node Gets Complete Copy of Graph
 - Every node “floods” network with data about its outgoing links
- Every Node Computes Routes to Every Other Node
 - Using single-source, shortest-path algorithm
- Every Node Updates Own Routing Table
- Process Performed Whenever Needed
 - When connections die / reappear
 - Periodically


Link State Pros & Cons

- Advantages
 - Rapidly adapts to changes in network
 - Can afford to use more sophisticated link costs (“metrics”)
 - Can incorporate multiple paths
- Disadvantages
 - Difficult to make reliable
 - Must make sure all nodes get consistent copy of graph
 - Tempting to add lots of features
 - Sources of complexity & bugs

Dijkstra's Algorithm

- Edsgar Dijkstra (1930--2002)
 - Pioneer in understanding mathematical basis for computer science
 - Fundamental ideas in concurrency (e.g., semaphores)
- Given
 - Graph with source node s and edge costs $c(u,v)$
 - Determine least cost path from s to every node v
- Shortest Path First Algorithm
 - Traverse graph in order of least cost from source

Dijkstra's Algorithm: Concept



• Node Sets

– Done

- Already have least cost path to it

– Horizon:

- Reachable in 1 hop from node in Done

– Unseen:

- Cannot reach directly from node in Done

• Label


– $d(v)$ = path cost

- From s to v
- All but v are done nodes
- Otherwise optimal

• Path


- Keep track of last link in path

Dijkstra's Algorithm: Initially




- No nodes done
- Source in horizon

Dijkstra's Algorithm: Selection




- Select node v in horizon with minimum $d(v)$

Dijkstra's Algorithm: Selection



- Add selected node to Done

Dijkstra's Algorithm: Update




- Update costs based on paths having last link from newly added Done node
 - Could change values of nodes in horizon
 - Could add new nodes to horizon
- Update link information

OSPF Routing Protocol

- Open
 - Open standard created by IETF
- Shortest-Path First
 - Another name for Dijkstra's algorithm
- Most Prevalent Intradomain Routing Protocol

Sending Link States by Flooding

- X Wants to Send Information
 - Sends on all outgoing links
- When Node Y Receives Information from Z
 - Send on all links other than Z
- Naïve Approach:
 - Floods indefinitely.
 - Prevent through sequence numbers



OSPF Reliable Flooding

- Transmit Link State Advertisements
 - Originating Router
 - List of directly connected neighbors of that node with the cost of the link to each one
 - Sequence Number
 - Incremented each time sending new link information
 - Link State Age
 - Packet expires when a threshold is reached,

OSPF Flooding Operation

- Node X Receives LSA from Node Y
 - With Sequence Number q
 - Looks for entry with same origin/link ID
- Cases
 - No entry present
 - Add entry, propagate to all neighbors other than Y
 - Entry present with sequence number $p < q$
 - Update entry, propagate to all neighbors other than Y
 - Entry present with sequence number $p > q$
 - Send entry back to Y
 - To tell Y that it has out-of-date information
 - Entry present with sequence number $p = q$
 - Ignore it

Flooding Issues


- When Should it be Performed
 - Periodically
 - When status of link changes
 - Detected by connected node
- What Happens when Router Goes Down & Back Up
 - Sequence number reset to 0
 - Other routers may have entries with higher sequence numbers
 - Router will send out LSAs with number 0
 - Will get back LSAs with last valid sequence number p
 - Router sets sequence number to $p+1$ & resends

Flooding Issues (Cont.)

- What if Sequence Number Wraps Around

- OSPF V1:

- Restrict LSAs to same semi-circle by regulating generation
 - But difficult to enforce with data corruption



- OSPF V2:

- Linear rather than circular space
 - Once sequence number reaches maximum, reset count to min
 - Flush out old sequence number by advertising LSA with MAXAGE
 - With 32-bit counter, doesn't happen very often

OSPF – Load Balancing

- Extensions that allow multiple paths to be taken for the same destination when they have equal costs
 - “ECMP” – Equal Cost Multi-Path.
- Key issue: packet reordering
 - Hash on IP header to ensure packets of the same flow take the same path.
- Recent body of research
 - ECMP does not work well
 - E.g., Multiple large flows hashed to same path and might cause congestion.
 - New solutions based on SDNs.


OSPF – Load Balancing

- Extensions that allow multiple paths to be taken for the same destination when they have equal costs
 - “ECMP” – Equal Cost Multi-Path.
- Key issue: packet reordering
 - Hash on IP header to ensure packets of the same flow take the same path.
- Recent body of research
 - ECMP does not work well
 - E.g., Multiple large flows hashed to same path and might cause congestion.
 - New solutions based on SDNs.

OSPF: Richer cost metrics

- How do we set costs?
 - Common recommendation: inversely proportional to link capacities
- Static cost or latency metrics not sufficient
 - Handling network congestion important
 - Traffic patterns might change over time
- Research in 2000's: (e.g., AT&T deployments)
 - Dynamic setting of weights in a manner that takes traffic demands into account
 - Changing weight settings: slower time-scale operation
- SDNs offer new opportunities.

OSPF Routing Hierarchy



- Partition Network into “Areas”
 - Router maintains link states of nodes within its area
 - Nodes in lower-level area use area-border router as default router
 - Backbone nodes can “summarize” routes within area

OSPF External Routes



- Limited Connectivity to Rest of Internet
 - Stub AS
 - Single border router
 - Can use border router as default for all addresses outside AS
 - Non-stub AS
 - External addresses need to be routed to appropriate border router
 - Can often summarize set of addresses by giving CIDR address

Emerging trends

- RIP Viewed as Outmoded
 - Good when networks small and routers had limited memory & computational power
 - Does not handle large networks with complex features
- OSPF Advantages
 - Fast convergence when configuration changes
 - Able to use more sophisticated metrics
- New trends
 - OSPF viewed overly complicated for more sophisticated traffic engineering tasks (load balancing, adjusting to congestion etc.)
 - Interest in SDNs to revisit routing